

喧

原

Aggregation Apparitions

Level 5 - Section 1

Common Aggregations

Finding the Number of Potions Per Vendor

Time for an audit! We need to know how many potions we have per vendor.



Name:

"Shrinking"
Vendor:

"Kettlecooked"
Price: 9.99

Name:

"Invisibility"
Vendor:

"Kettlecooked"
Price: 15.99
...

Name:

"Luck"
Vendor:

"Leprechau..."
Price: 59.99

```
(!)
```

```
> db.potions.find({}, {"name": true, "vendor": true})
```

We could manually pull all the data and count everything, but it's better to have MongoDB handle that for us!



Introducing the Aggregation Framework

The aggregation framework allows for advanced computations.

Takes stage operators as parameters

```
> db.potions.aggregate(
  [{"$group": {"_id": "$vendor_id"}}]
)
```

Go within an array

Stage operator that's used to group data by any field we specify

Field names that begin with a "\$" are called "field paths" and are links to a field in a document



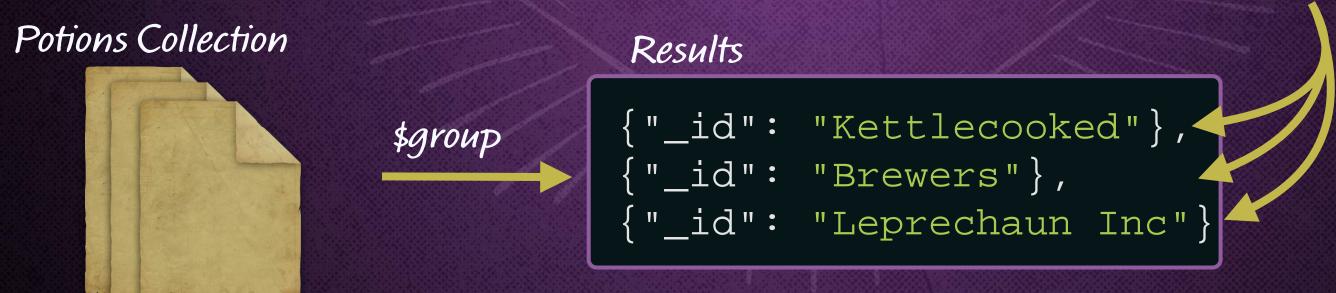
"Aggregate" is a fancy word for combining data.

Using the Aggregation Framework to Group Data

```
> db.potions.aggregate(
  [{"$group": {"_id": "$vendor_id"}}]
)
```

This is known as the "group key" and is required

Returns result object containing the unique vendors in the inventory collection



Using Accumulators

Anything specified after the group key is considered an accumulator. Accumulators take a single expression and compute the expression for grouped documents.

```
SHELL
> db.potions.aggregate([
                                        "total": { "$sum": 1}}
 { "$group": _{ "_id": "$vendor_id",
                                     Will add 1 for each
                                                                  Accumulator
            Group key
                                    matching document
 Inventory
                                   Total number of documents per vendor!
                         Results
                          "_id": "Kettlecooked", "total": 2},
                $group
                          "_id": "Brewers", "total": 1,},
                          {"_id": "Leprechaun Inc", "total": 1
```

Field Paths Vs. Operators

When values begin with a "\$", they represent field paths that point to the value

When fields begin with a "\$", they are operators that perform a task

Summing the Grade Per Vendor

Results

```
{"_id": "Kettlecooked", "total": 2, "grade_total": 400},

{"_id": "Brewers", "total: 1, "grade_total": 340},

{"_id": "Leprechaun Inc", "total": 1, "grade_total": 92}
```

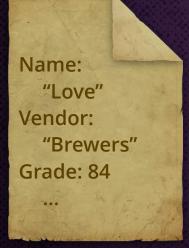
Averaging Potion Grade Per Vendor



Name:

"Shrinking"
Vendor:

"Kettlecooked"
Grade: 94



Name:

"Sleep"
Vendor:

"Brewers"
Grade: 30

Results

```
{"_id": "Kettlecooked", "avg_grade": 82 },
{"_id": "Brewers", "avg_grade": 57 }
```

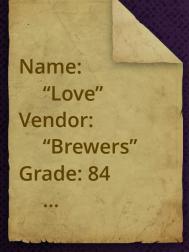
Returning the Max Grade Per Vendor



Name:

"Shrinking"
Vendor:

"Kettlecooked"
Grade: 94



Name:

"Sleep"
Vendor:

"Brewers"
Grade: 30

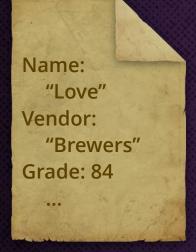
Results

```
{"_id": "Kettlecooked", "max_grade": 94},
{"_id": "Brewers", "max_grade": 84}
```

Using \$max and \$min Together



Name: "Shrinking" Vendor: "Kettlecooked" Grade: 94



Name: "Sleep" Vendor: "Brewers" Grade: 30

```
SHELL
> db.potions.aggregate([
 { "$group":
    "_id": "$vendor_id",
    "max_grade": { "$max": "$grade" } ,
    "min_grade": { "$min": "$grade" }
```

Results

We can use the same field in

multiple accumulators

```
"_id": "Kettlecooked", "max_grade": 94, "min_grade": 70 },
"_id": "Brewers", "max_grade": 84, "min_grade": 30 }
```

Aggregation Apparitions

Level 5 – Section 2

The Aggregation Pipeline

Pulling Conditional Vendor Information

It turns out that potions made with unicorn aren't permitted, so we need to count the number of potions per vendor that contain it.

Notice

All potions containing unicorn are strictly forbidden

Steps to find potions with unicorns:

- 1) Query potions
- 2) Group by vendor
- 3) Sum the number of potions per vendor



Introducing the Aggregation Pipeline

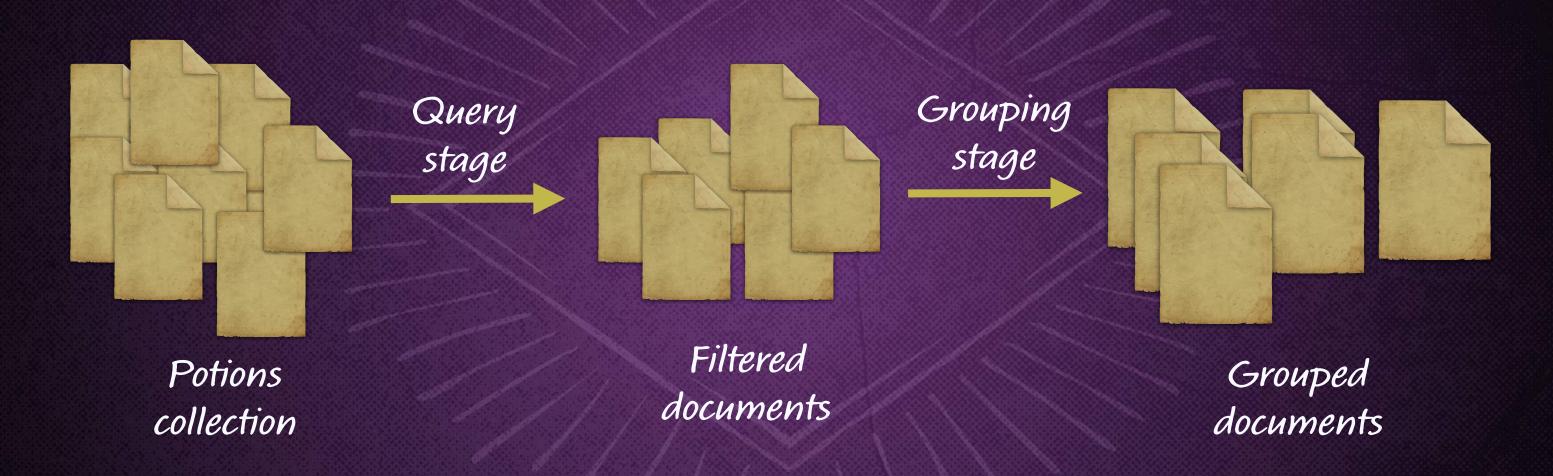
The aggregate method acts like a pipeline, where we can pass data through many stages in order to change it along the way.

```
db.potions.aggregate([stage, stage, stage])
```

We've already seen what a stage looks like in the last section

How the Pipeline Works

Each stage modifies the working data set and then passes the altered documents to the next stage until we get our desired result.



Using the \$match Stage Operator

\$match is just like a normal query and will only pass documents to the next stage if they meet the specified condition(s).

We can use the same query we would use with find()



Use match early to reduce the number of documents for better performance.

Grouping Potions With Data

```
SHELL
db.potions.aggregate([
 { "$match": { "ingredients": "unicorn" } } ,
 { "$group":
                                                2 stages separated by
      "_id": "$vendor_id",
                                               comma within an array
     "potion_count": { "$sum": 1}
```

Result

```
{"_id": "Poof", "potion_count": 20},
{"_id": "Kettlecooked", "potion_count": 1}
```

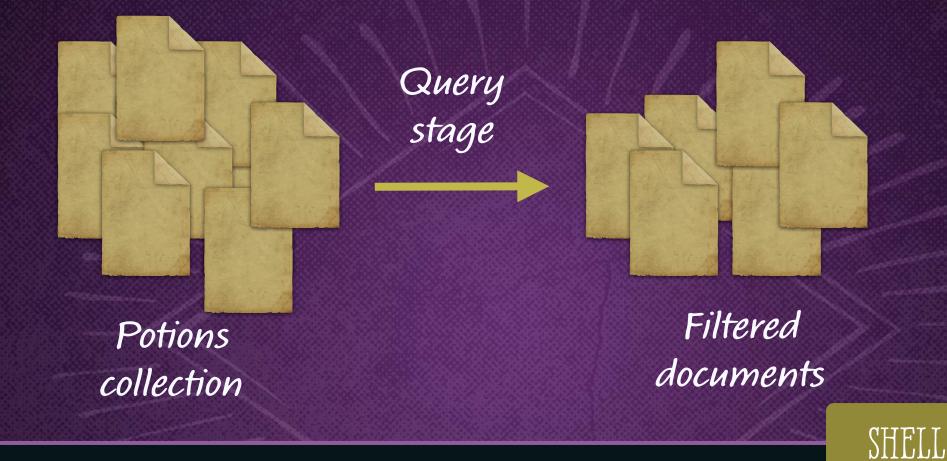
Top 3 Vendors With Potions Under \$15

Best value

- 1) Query for potions with a price less than 15
- 2) Group potions by vendor and average their grades
- 3) Sort the results by grade average
- 4) Limit results to only 3 vendors



Matching Potions Under \$15



```
db.potions.aggregate([
    {"$match": {"price": {"$lt": 15}}}
])
```

It's good practice to limit the number of results early on

Matching Potions Under \$15

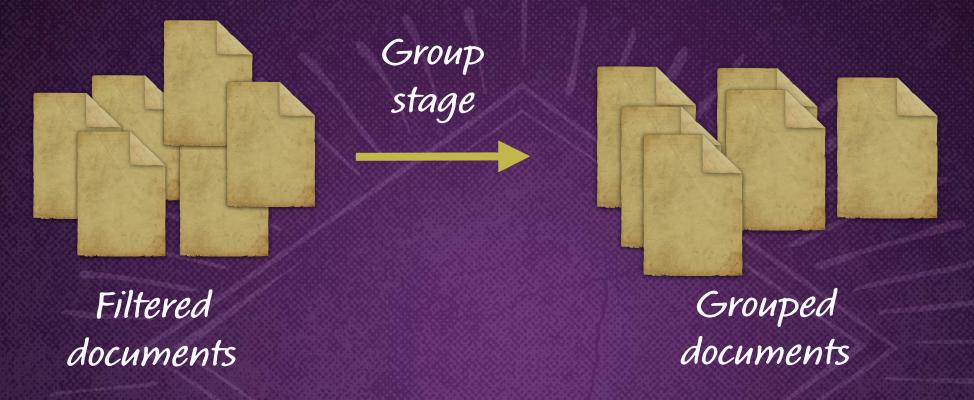


Filtered documents

```
SHELL
```

```
db.potions.aggregate([
    {"$match": {"price": {"$lt": 15}}}
])
```

Grouping Potions by Vendor



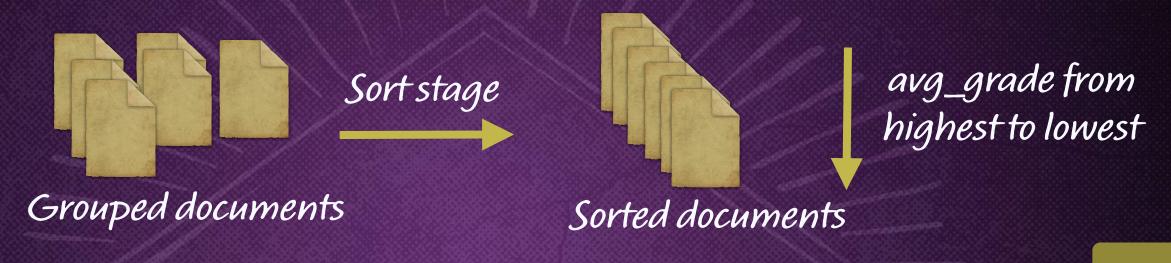
```
db.potions.aggregate([
    {"$match": {"price": {"$lt": 15}}},
    {"$group": {"_id": "$vendor_id", "avg_grade": {"$avg": "$grade"}}}
])
```

Vendor name as group key

Average of potion grades

Sorting Vendors by Average Grade

We can sort potions by using the **\$sort** stage operator.

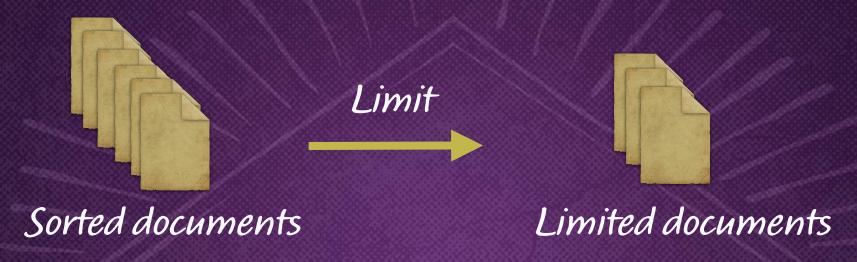


We can sort the field we created during the group stage

Sort by the field in descending order

Limiting the Number of Documents

We can use the *\$limit* stage operator to limit the number of documents that get passed on.



Optimizing the Pipeline

It's best practice to only send the needed data from stage to stage. Let's see what data we really need.

Only need vendor and grade for each potion after the match stage

Projections While Aggregating

We can limit the fields we send over by using **\$project**, which functions the same way as projections when we're querying with **find()**.

Vendor and grade for each potion after the match stage

We want to use \$project as soon as possible



It's common to see \$match and \$project used together early on and throughout the pipeline.

Aggregation Results

```
{"_id": "Kettlecooked", "avg_grade": 99 },
{"_id": "Leprechaun Inc", "avg_grade": 95 },
{"_id": "Brewers", "avg_grade": 90 }
```

Group key

Average potion grade